

# Server in Java

Adam Lider

# Contents

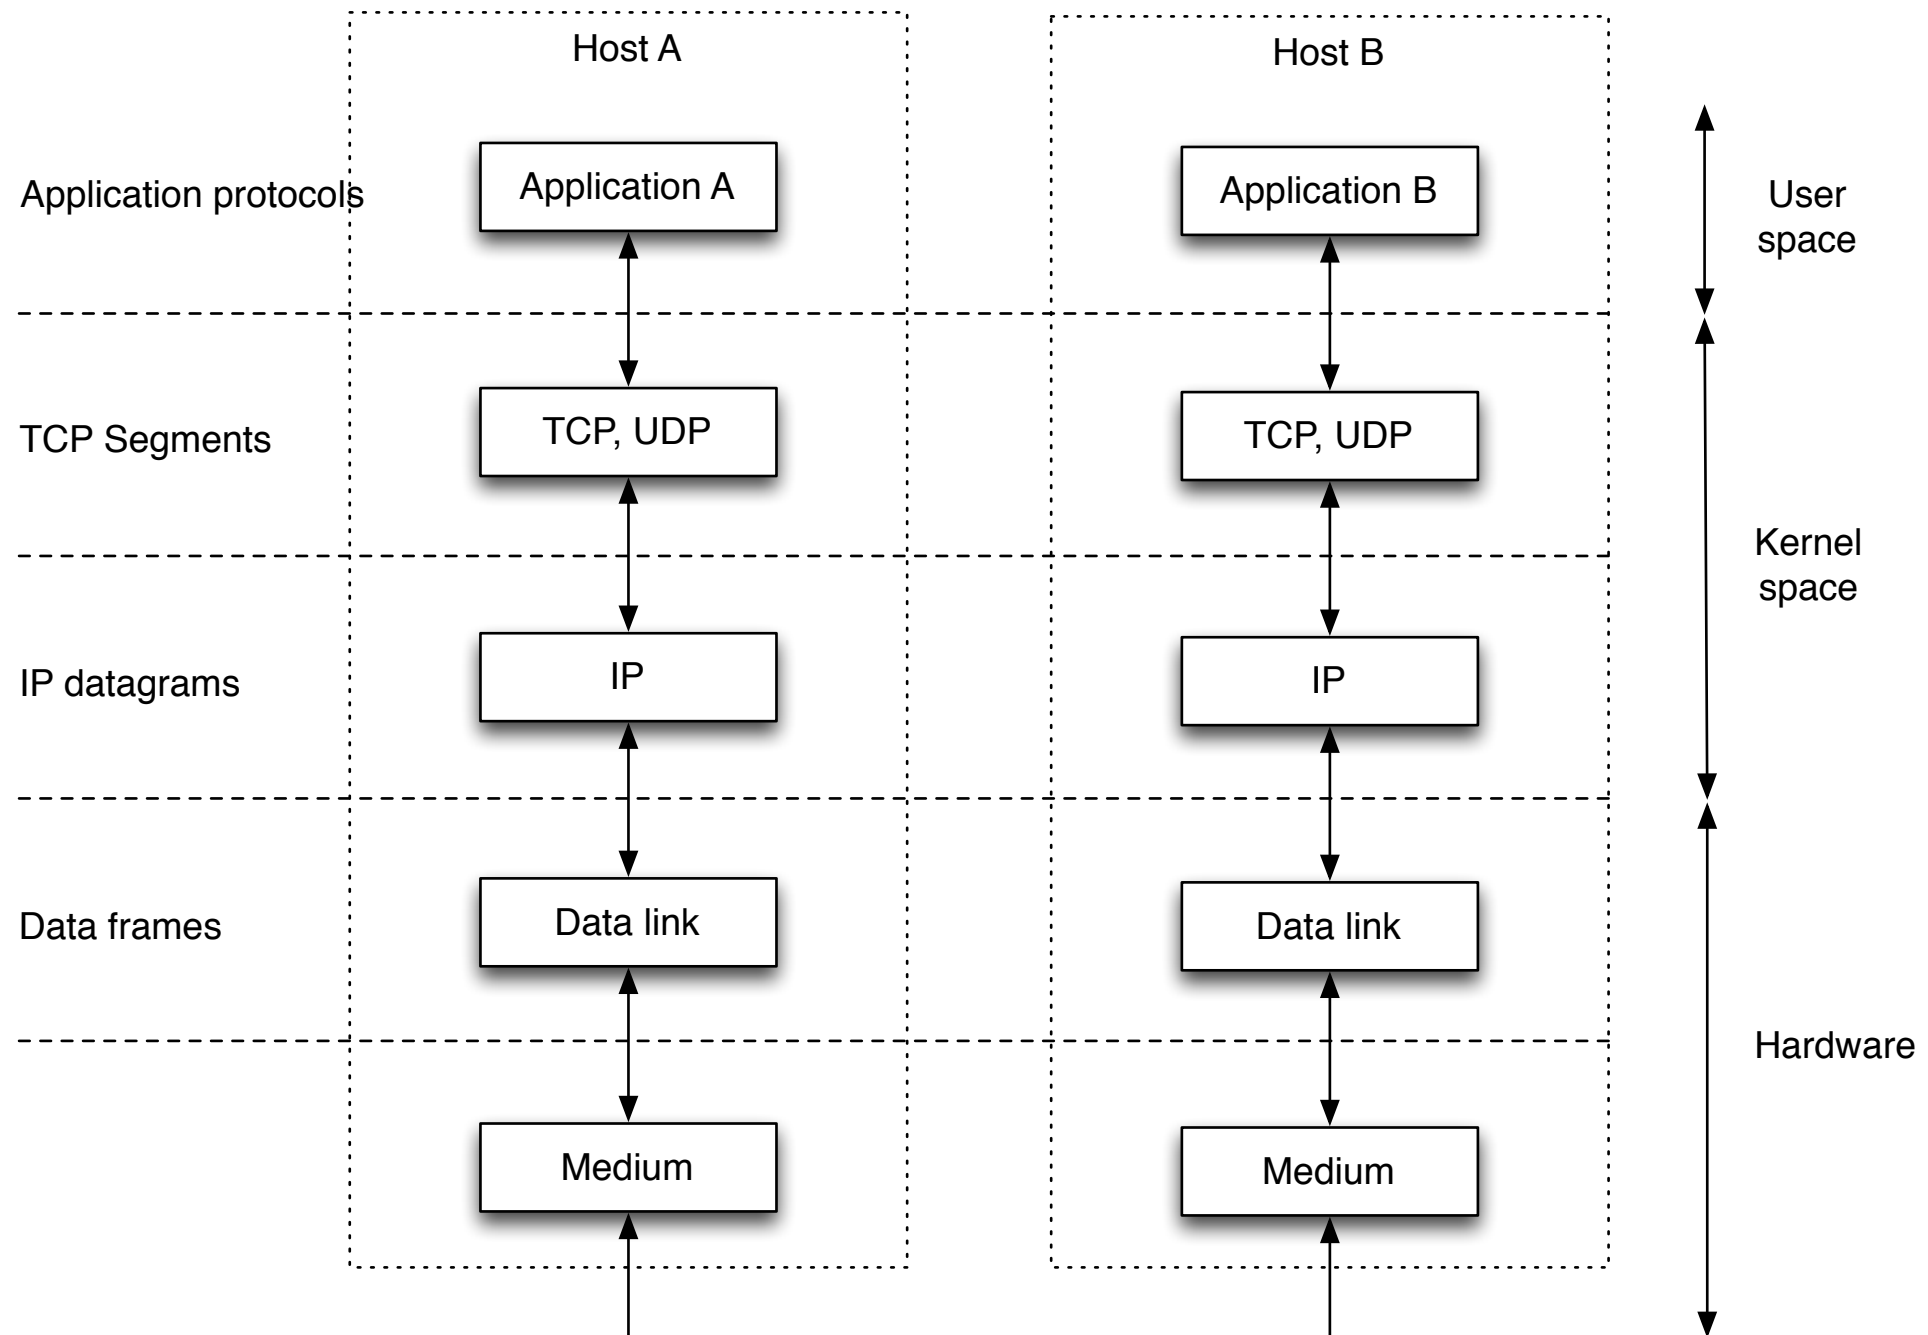
- Basics
- Custom protocol
- TCP details
- IO models
- Network architectures patterns

# Basics

- Layered communication
- Establishing connections
- Exchanging data
- Closing connections

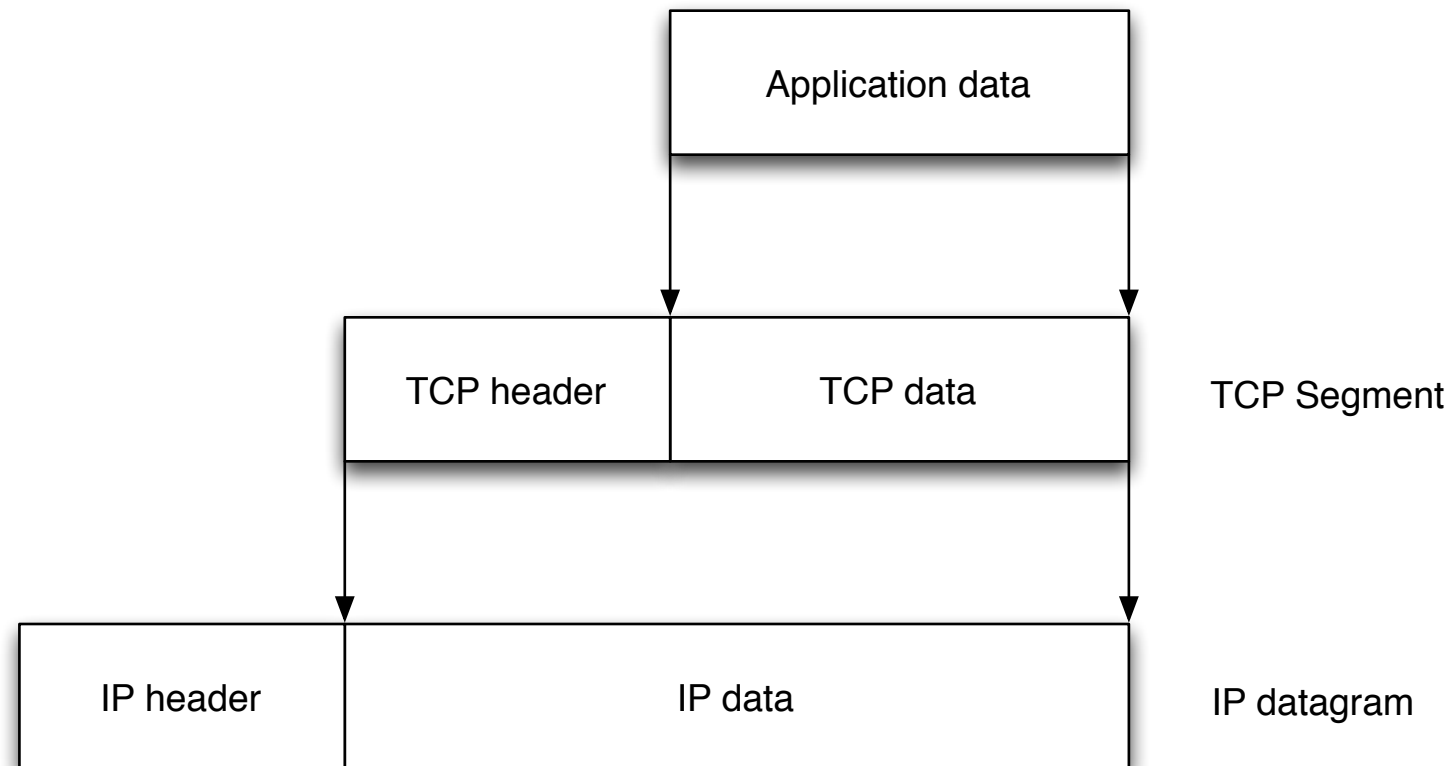
# Basics

## Layered communication



# Basics

## Layered communication - encapsulation



# Basics

## Establishing connections

### Client

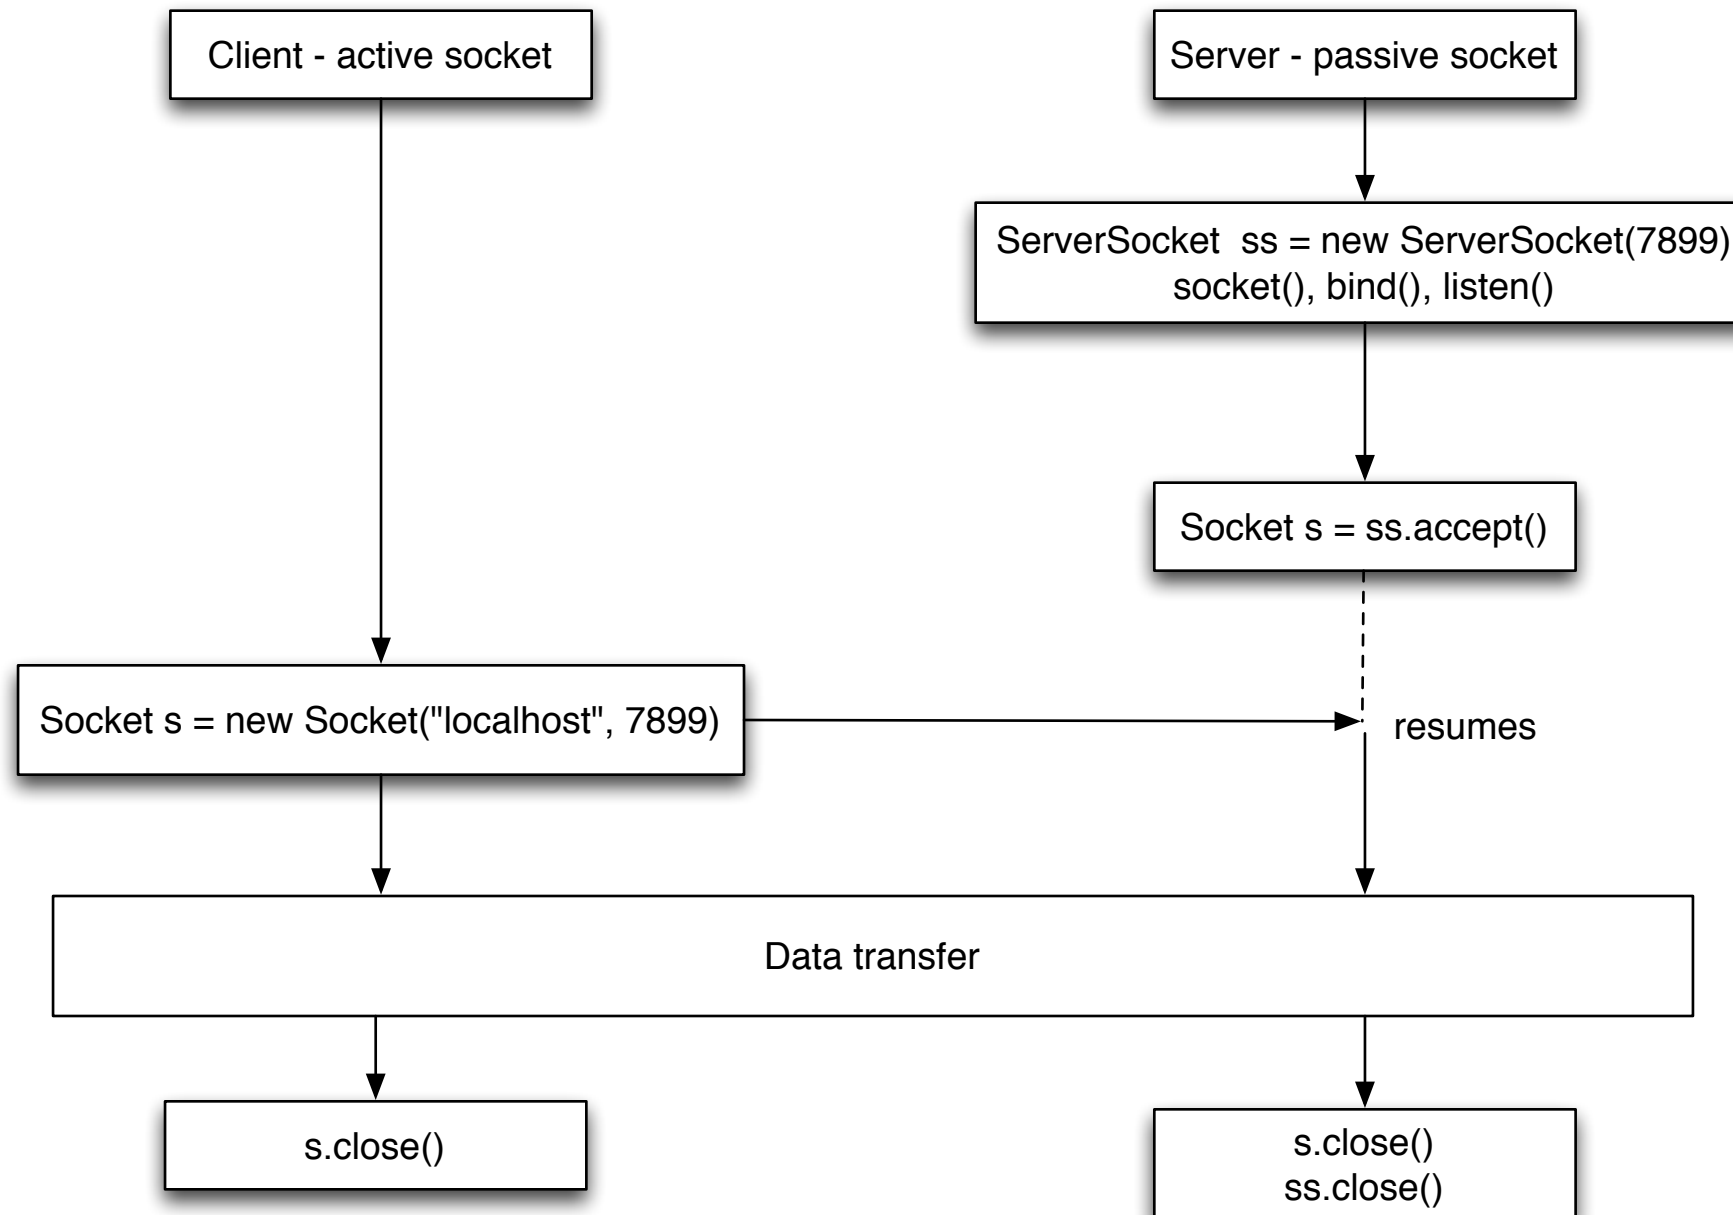
```
Socket s = new Socket("localhost", 7899);  
OutputStream out = s.getOutputStream();  
out.write(new byte[1024]);  
s.close();
```

### Server

```
ServerSocket ss = new ServerSocket(7899);  
Socket s = ss.accept();  
InputStream in = s.getInputStream();  
byte[] buffer = new byte[1024 * 8];  
int read = in.read(buffer);  
s.close();  
ss.close();
```

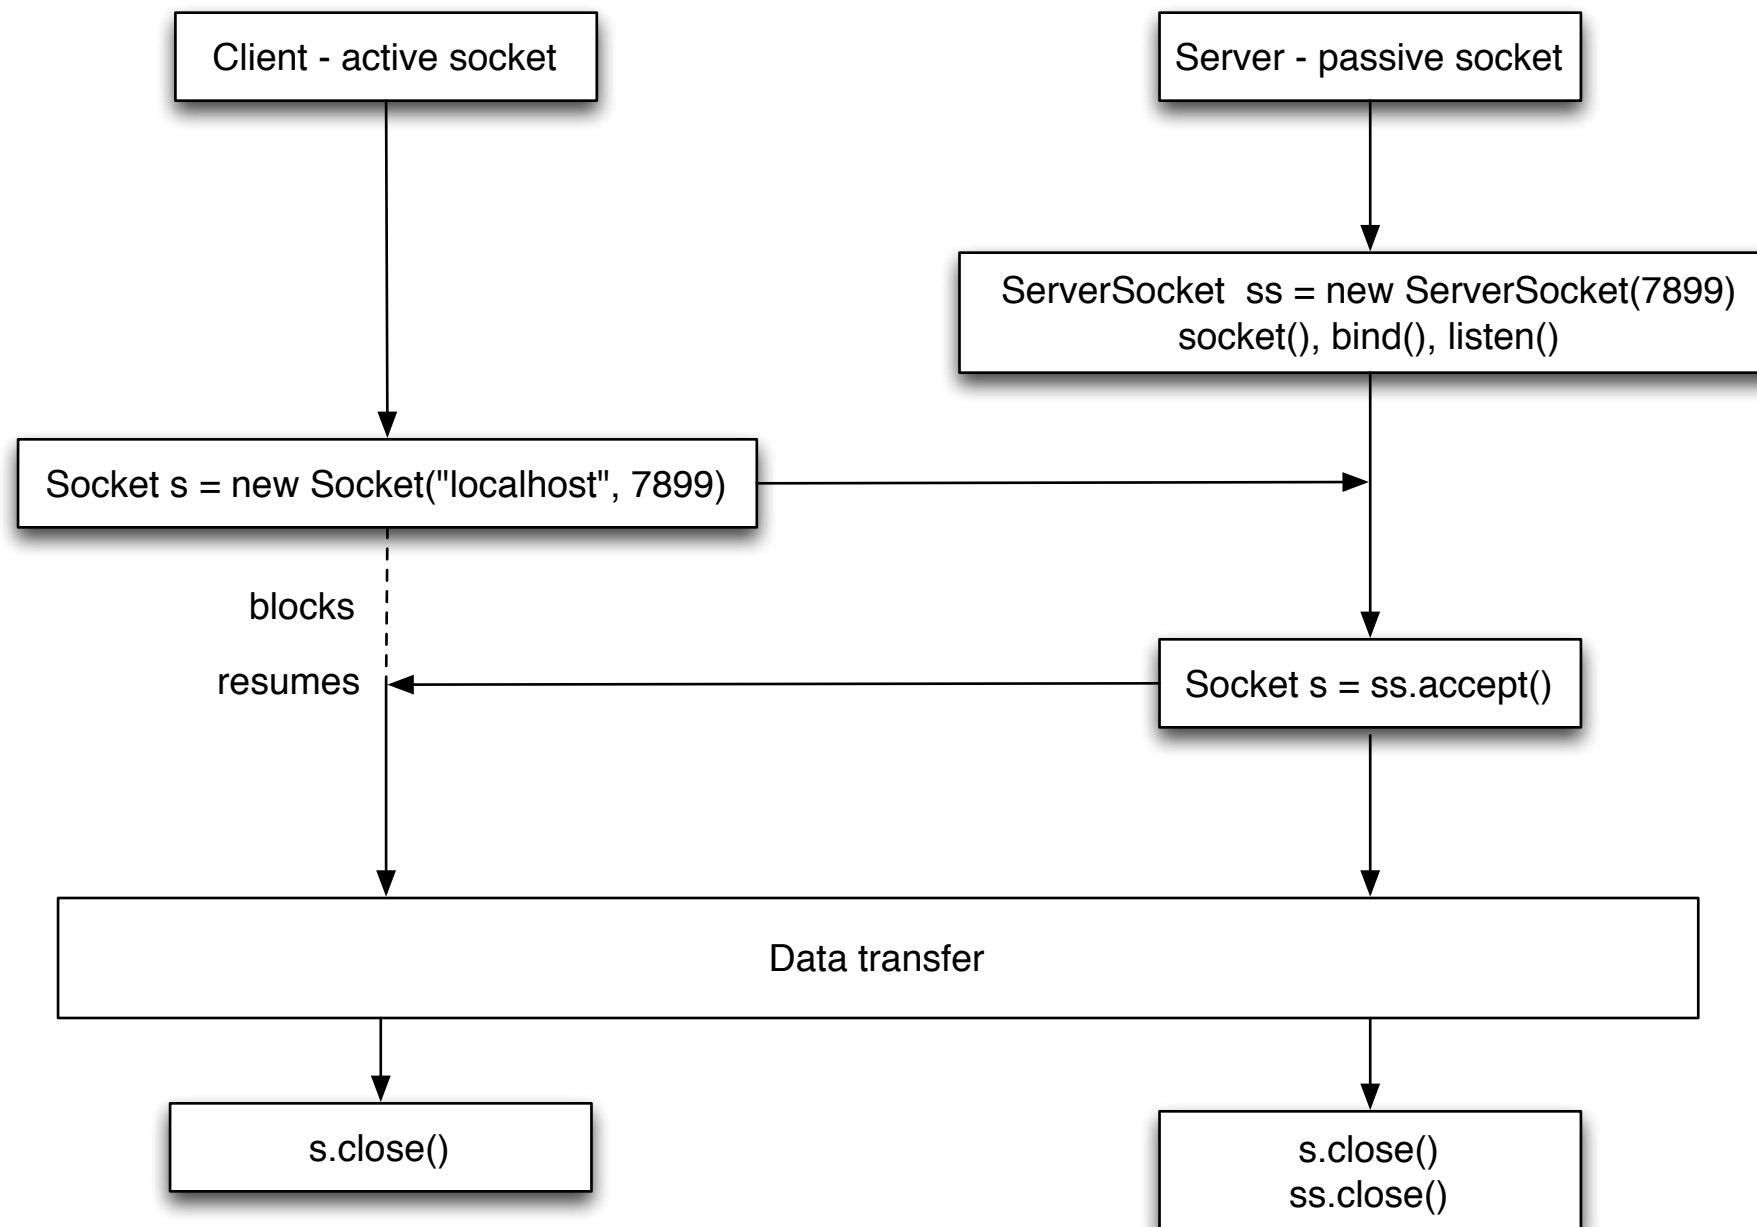
# Basics

## Establishing connections



# Basics

## Establishing connections - backlog



Java: default 50, `new ServerSocket(7899, 400);`  
Linux: default 128, `SOMAXCONN`



# Basics


## Establishing connections - three-way handshake

Client

```
Socket s = new Socket("localhost", 7899);  
OutputStream out = s.getOutputStream();  
out.write(new byte[1024]);  
s.close();
```

Server

```
ServerSocket ss = new ServerSocket(7899);  
Socket s = ss.accept();  
InputStream in = s.getInputStream();  
byte[] buffer = new byte[1024 * 8];  
int read = in.read(buffer);  
s.close();  
ss.close();
```



# Basics

## Data transfer

### Client

```
Socket s = new Socket("localhost", 7899);  
OutputStream out = s.getOutputStream();  
out.write(new byte[1024]);  
s.close();
```

### Server

```
ServerSocket ss = new ServerSocket(7899);  
Socket s = ss.accept();  
InputStream in = s.getInputStream();  
byte[] buffer = new byte[1024 * 8];  
int read = in.read(buffer);  
s.close();  
ss.close();
```

# Basics

## Data transfer

### Client

```
Socket s = new Socket("localhost", 7899);  
OutputStream out = s.getOutputStream();  
out.write(new byte[1024]);  
out.write(new byte[1024 * 2]);  
out.write(new byte[1024 * 4]);  
s.close();
```

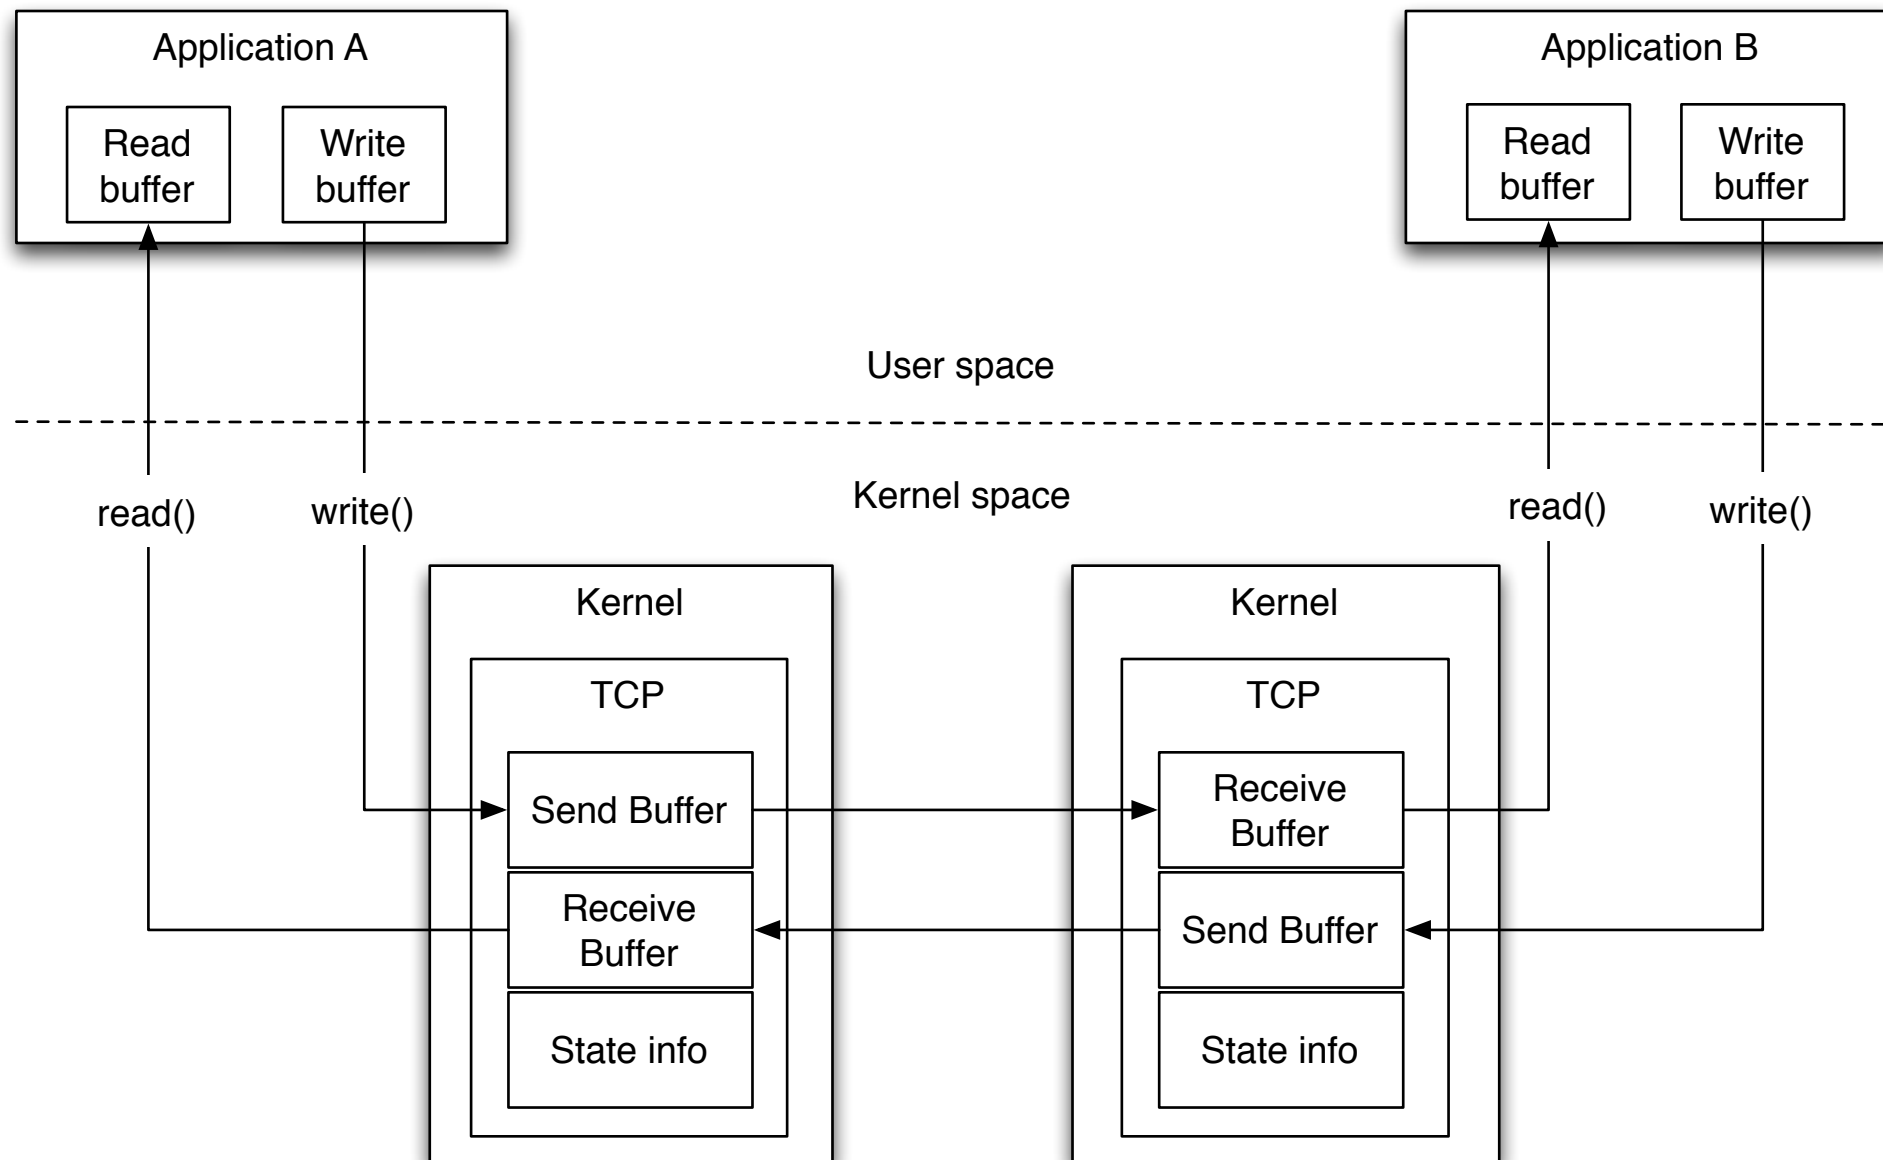
### Server

```
ServerSocket ss = new ServerSocket(7899);  
Socket s = ss.accept();  
InputStream in = s.getInputStream();  
byte[] buffer = new byte[1024 * 8];  
int read = in.read(buffer);  
s.close();  
ss.close();
```

There is no any correspondence between writes performed at one end of connection and reads at the other end.

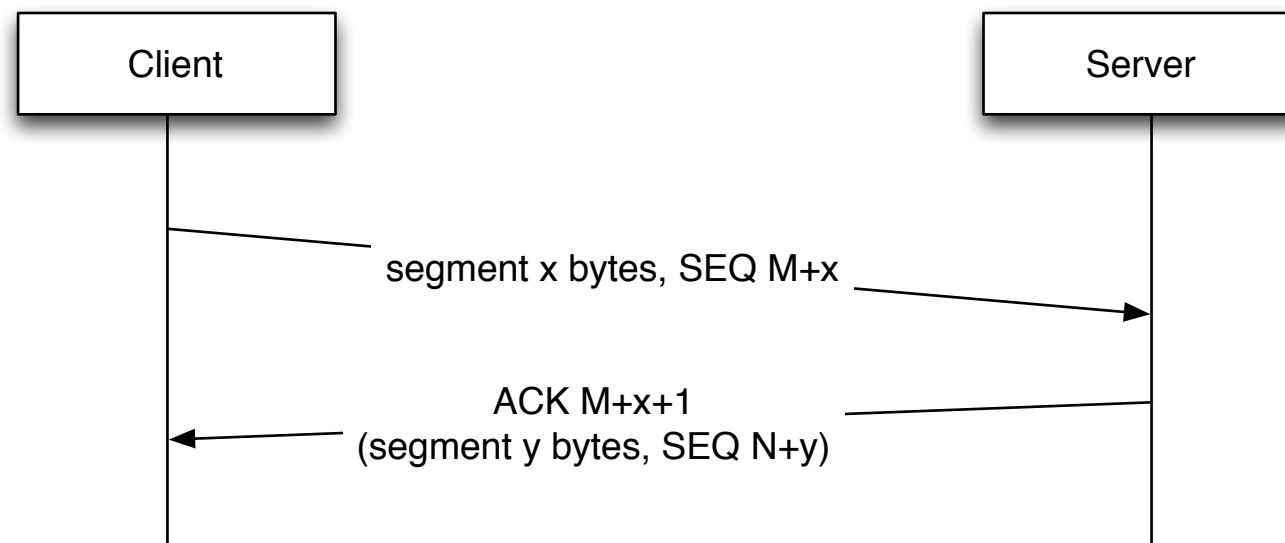
# Basics

## Data transfer



# Basics

## Data transfer



# Basics

## Closing connections

### Client

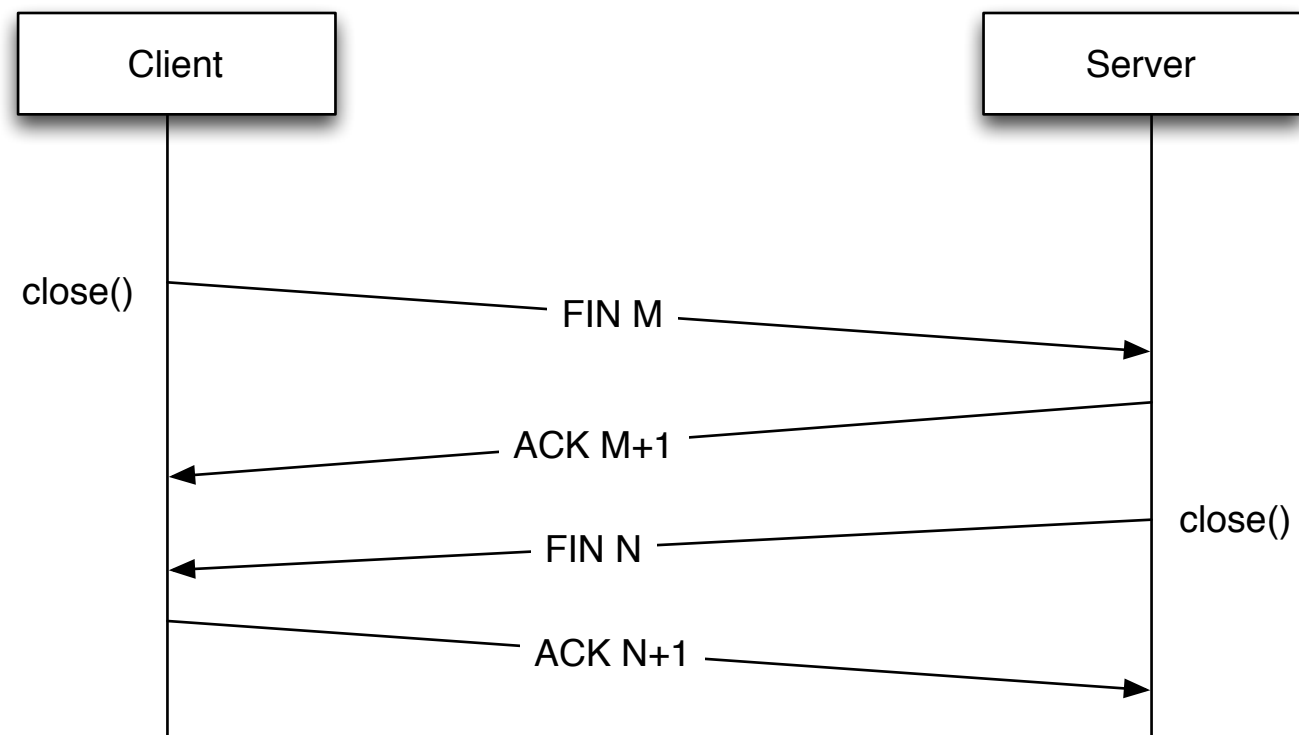
```
Socket s = new Socket("localhost", 7899);  
OutputStream out = s.getOutputStream();  
out.write(new byte[1024]);  
s.close();
```

### Server

```
ServerSocket ss = new ServerSocket(7899);  
Socket s = ss.accept();  
InputStream in = s.getInputStream();  
byte[] buffer = new byte[1024 * 8];  
int read = in.read(buffer);  
s.close();  
ss.close();
```

# Basics

## Closing connections



Custom protocol



# Custom protocol

- Encodings
- Framing messages

# Custom protocol

## Encodings

- Integers
  - big vs little endian
  - signed vs unsigned
- Strings of text

All parties must agree on the representation for integer and strings of text

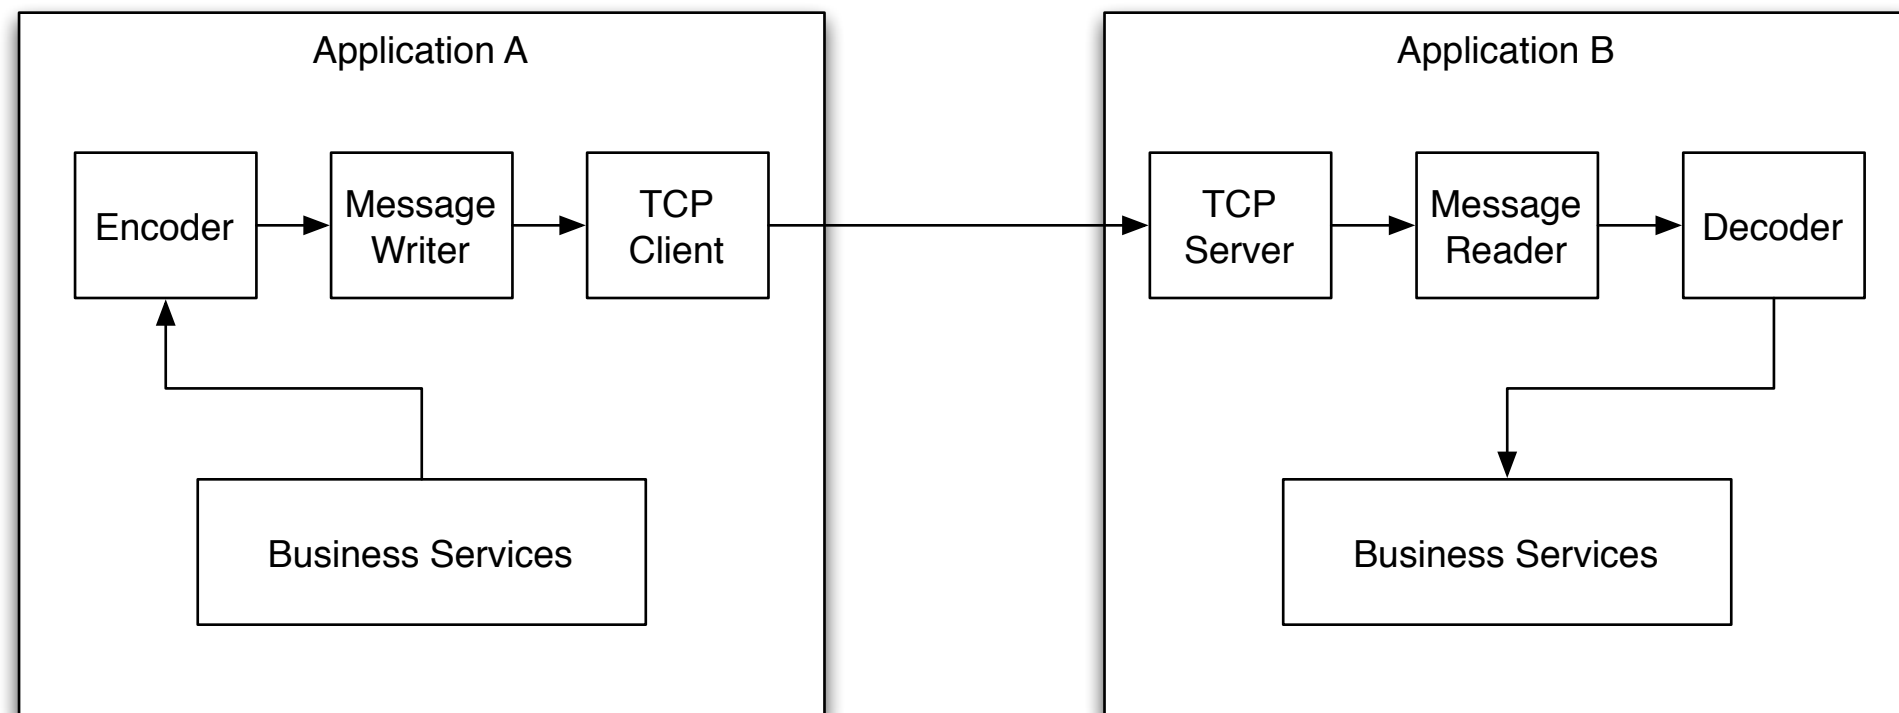
# Custom protocol

## Framing and parsing

- Separation of concerns
- Framing
  - Delimiter
  - Length

# Custom protocol

Separation of concerns - Framing and parsing



# TCP details

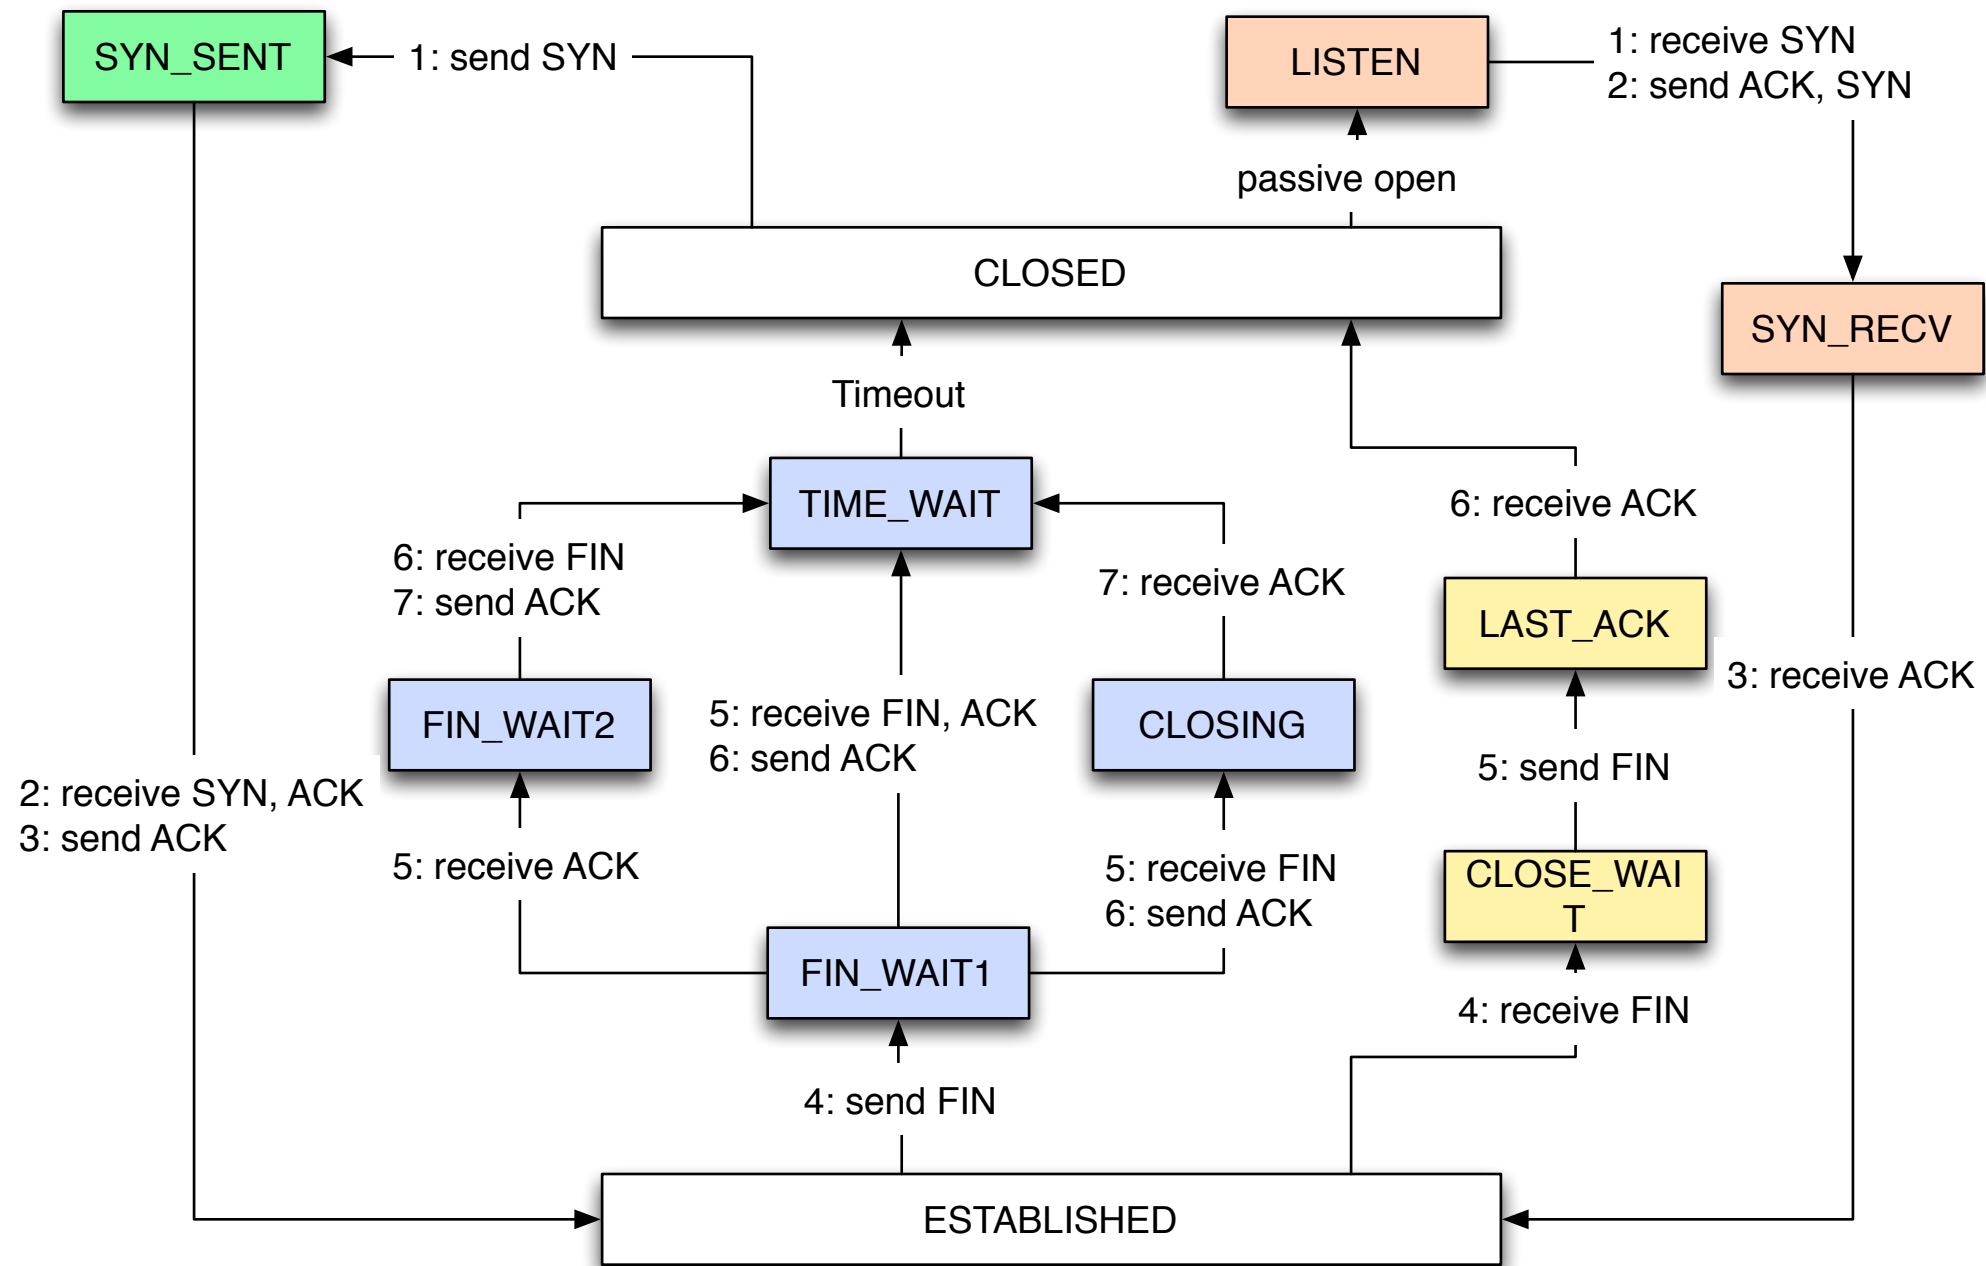
- State diagram
- Nagle's algorithm
- IP packet fragmentation
- Half-duplex connection termination
- Socket options

# TCP details

- Ordered data (byte) transfer
- Retransmission of lost packets
- Error-free data transfer
- Flow control
- Congestion control

# TCP details

## State diagram



# TCP details

## Nagle's algorithm

- If there is sufficient data to compromise an entire TCP segment (data  $\geq$  MSS) then send it immediately
- If there is no pending data in local buffers and no pending ACK of receipt from receiving end, then send it immediately
- If there is a pending ACK of receipt from receiving end and not enough data to compromise an entire TCP segment, then put the data into the local buffer



# TCP details

## IP packet fragmentation

- MSS - Maximum Segment Size
- MTU - Maximum transmission unit

# TCP details

## Half-duplex connection termination

### Client

```
Socket s = new Socket("localhost", 7899);
OutputStream out = s.getOutputStream();
out.write(new byte[1024]);
out.write(new byte[2048]);
out.write(new byte[4096]);
out.close();
s.close();
```

### Server

```
ServerSocket ss = new ServerSocket(7899);
Socket s = ss.accept();

OutputStream out = s.getOutputStream();
out.write("Hello".getBytes());

InputStream in = s.getInputStream();
byte[] buffer = new byte[1024 * 8];
int read = in.read(buffer);

s.close();
ss.close();
```

# TCP details

## Half-duplex connection termination

RFC 1122, section 4.2.2.13

“A host MAY implement a ‘half-duplex’ TCP close sequence, so that an application that has called CLOSE cannot continue to read data from the connection. If such a host issues a CLOSE call while received data is still pending in TCP, or if new data is received after CLOSE is called, its TCP SHOULD send a RST to show that data was lost.”

# TCP details

## shutdown output

### Client

```
Socket s = new Socket("localhost", 7899);

OutputStream out = s.getOutputStream();
out.write(new byte[1024]);
out.write(new byte[2048]);
out.write(new byte[4096]);
s.shutdownOutput();

InputStream in = s.getInputStream();
// read data

s.close();
```

### Server

```
ServerSocket ss = new ServerSocket(7899);
Socket s = ss.accept();

InputStream in = s.getInputStream();
// read data

OutputStream out = s.getOutputStream();
out.write("OK".getBytes());

s.close();
ss.close();
```

# TCP details

## Read and write deadlock

- Both ends read
- Both ends write

# TCP details

## Socket options

- TCP\_NODELAY, setTcpNoDelay
- SO\_REUSEADDR, setReuseAddress
- SO\_RCVTIMEO, setSoTimeout ?
- SO\_SNDTIMEO
- SO\_SNDBUF, setSendBufferSize
- SO\_RCVBUF, setReceiveBufferSize
- SO\_LINGER, setSoLinger

# IO models

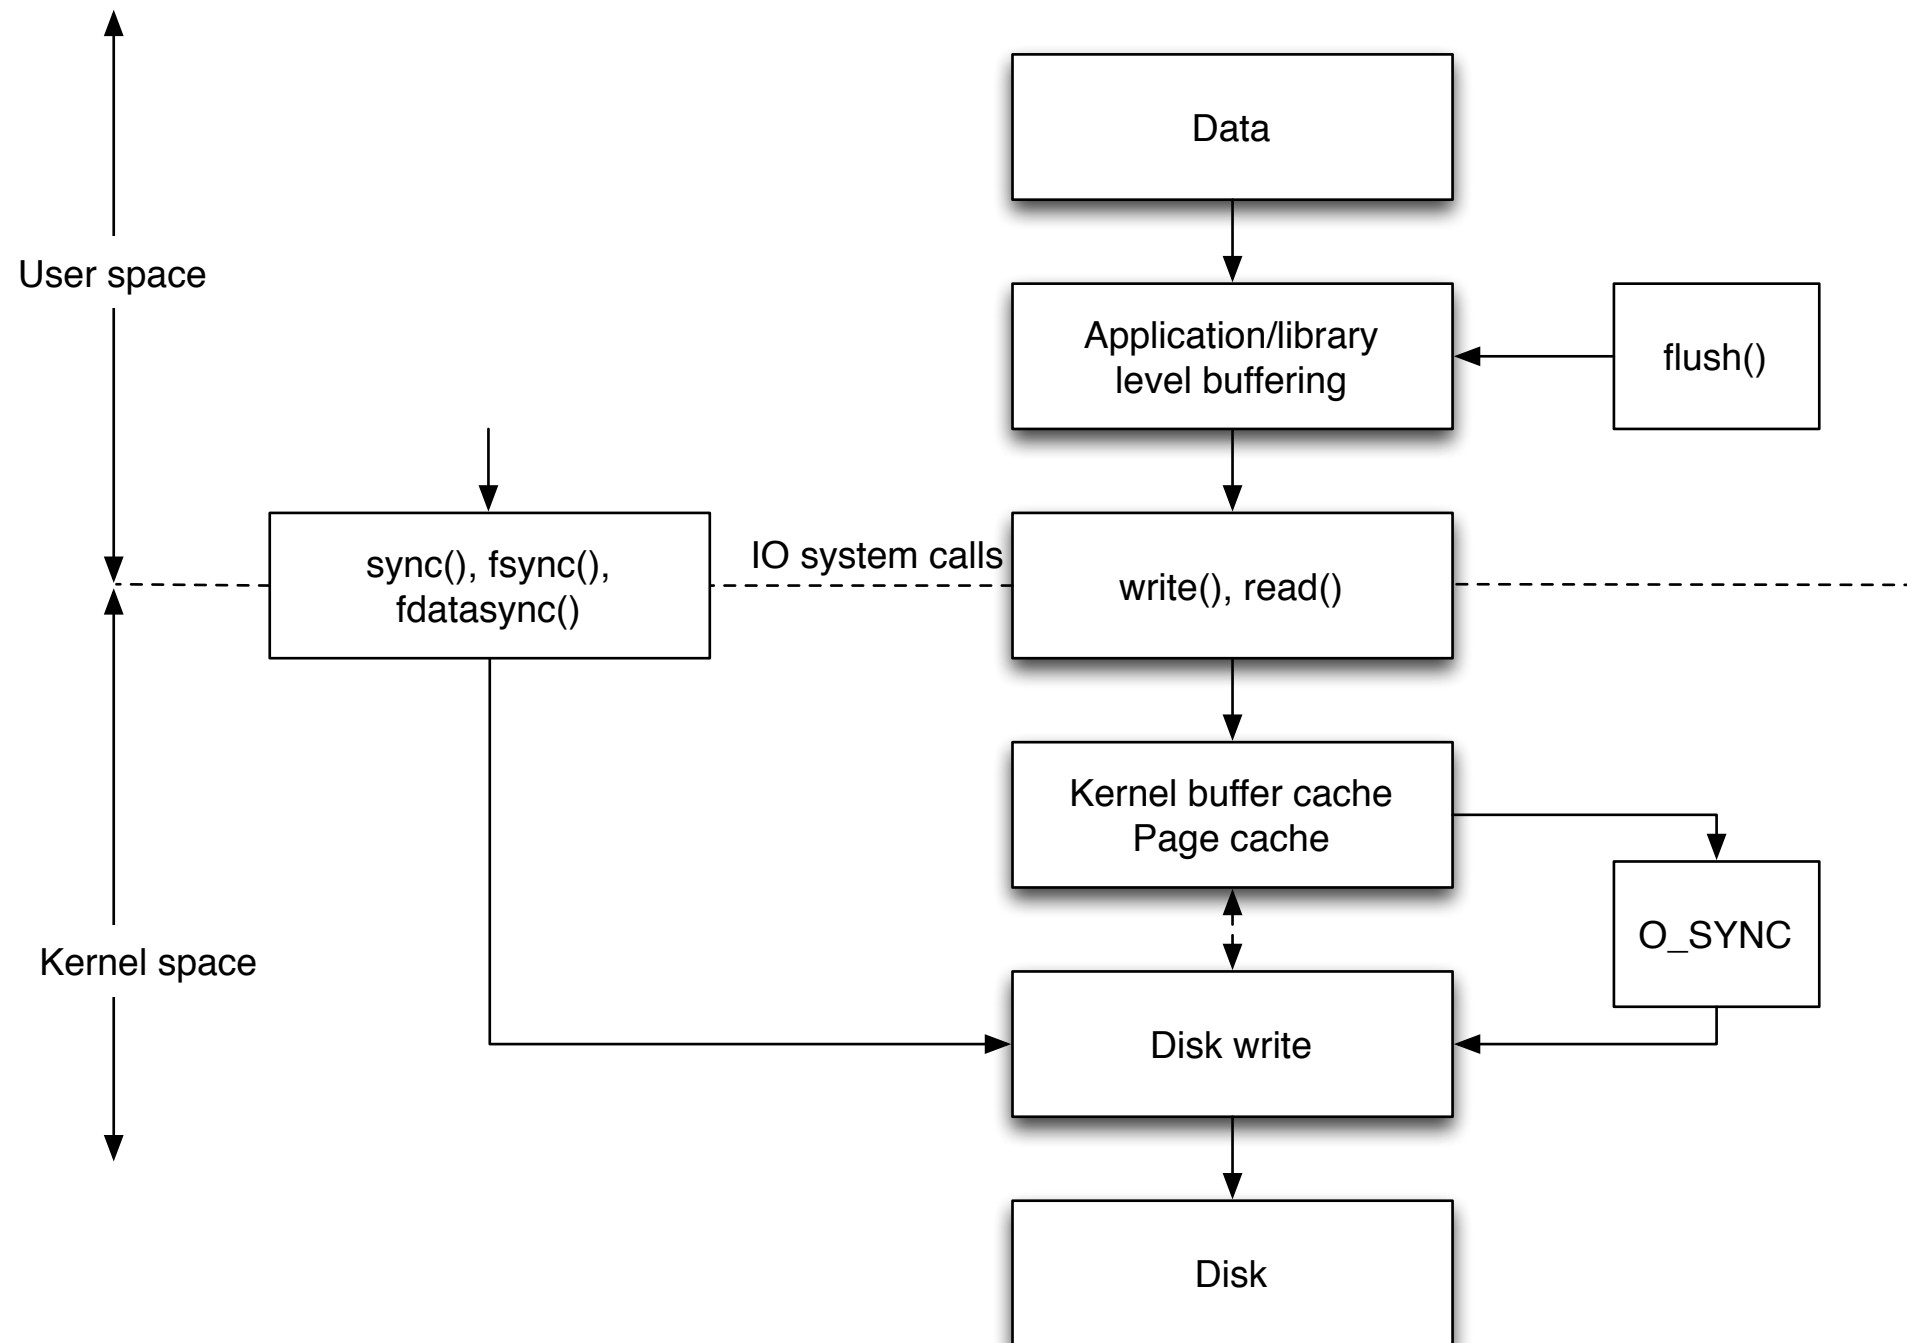
# IO models

- File IO
- Blocking, non-blocking, synchronous
- Java IO, NIO, NIO.2



# IO models

## File IO



# IO models

## Non-blocking

- Sockets, pipes, terminals, FIFOs
- O\_NONBLOCK
- Monitoring file descriptors:
  - select, poll
  - epoll
  - kqueue

# IO models

## Asynchronous

- libaio
- POSIX AIO
- O\_DIRECT

# IO models

## Java IO, NIO, NIO.2

- Java IO: Streams, blocking IO
- Java NIO:
  - Channels, Buffers
  - Blocking, non-blocking IO
- Java NIO.2: Asynchronous IO

# IO models

## Java NIO - non-blocking, sockets

```
RequestHandler handler = new RequestHandler();
long timeout = 1000;

Selector selector = Selector.open();

ServerSocketChannel listenChannel = ServerSocketChannel.open();
listenChannel.bind(new InetSocketAddress("localhost", 7789));
listenChannel.configureBlocking(false);
listenChannel.register(selector, SelectionKey.OP_ACCEPT);

for (;;) {
    if (selector.select(timeout) != 0) {
        continue;
    }
    Iterator<SelectionKey> selectedKeys = selector
                                                .selectedKeys()
                                                .iterator();

    while (selectedKeys.hasNext()) {
        SelectionKey key = (SelectionKey) selectedKeys.next();
        selectedKeys.remove();

        if (key.isAcceptable()) {
            handler.handleAccept(key);
        }
        if (key.isReadable()) {
            handler.handleRead(key);
        }
        if (key.isValid() && key.isWritable()) {
            handler.handleWrite(key);
        }
    }
}
```

# IO models

## Java NIO - non-blocking, sockets

```
public static class RequestHandler {

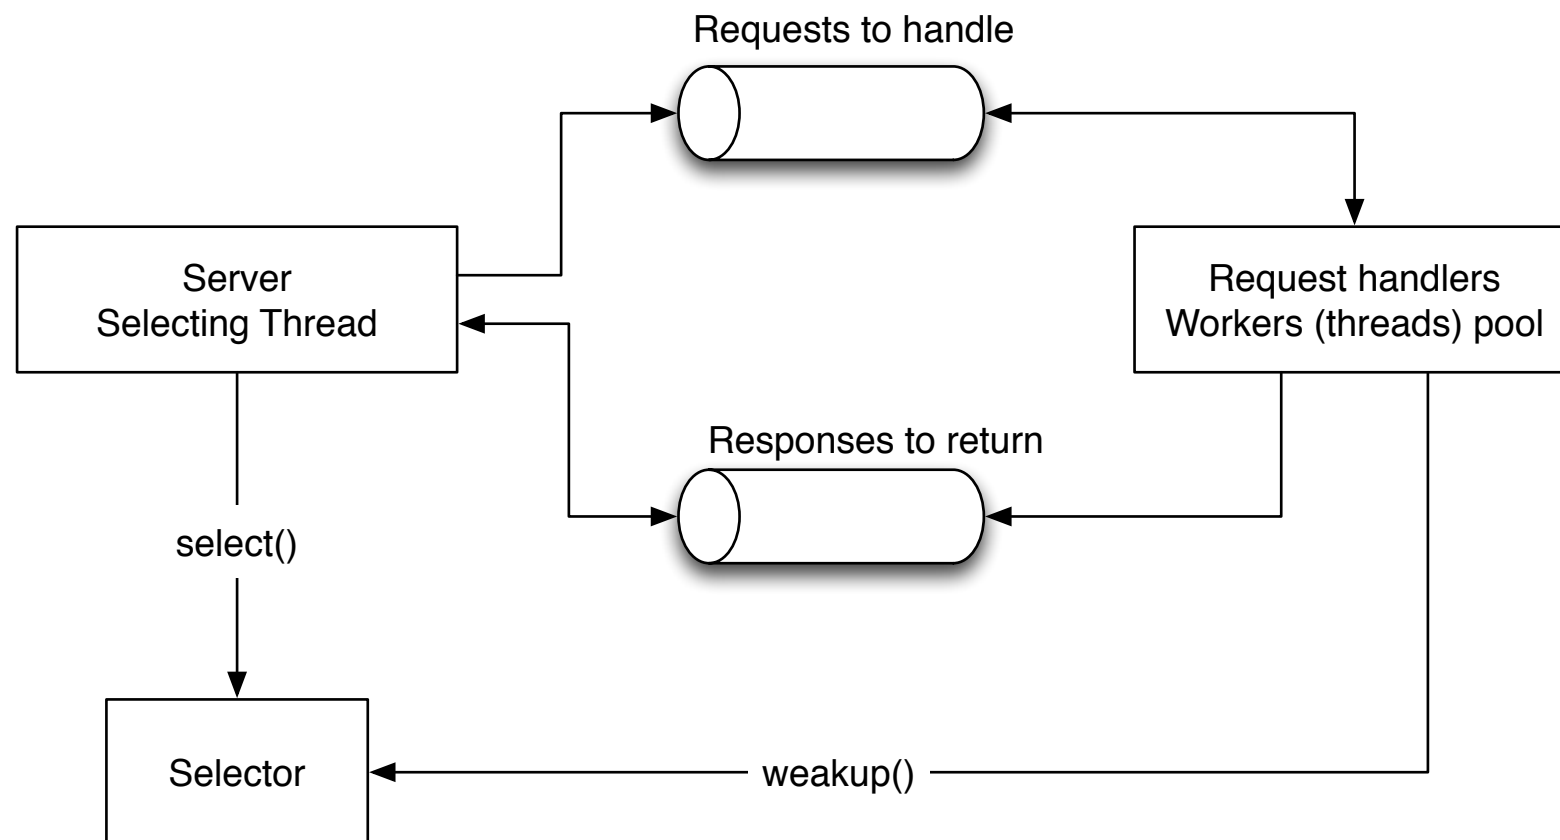
    public void handleAccept(SelectionKey key) throws IOException {
        ServerSocketChannel serverChannel = (ServerSocketChannel) key.channel();
        SocketChannel clientChannel = serverChannel.accept();
        clientChannel.configureBlocking(false);
        clientChannel.register(
            key.selector(),
            SelectionKey.OP_READ,
            ByteBuffer.allocate(1024));
    }

    public void handleRead(SelectionKey key) throws IOException {
        SocketChannel channel = (SocketChannel) key.channel();
        ByteBuffer buffer = (ByteBuffer) key.attachment();
        int read = channel.read(buffer);
        if (read == -1) {
            channel.close();
        } else {
            key.interestOps(SelectionKey.OP_READ | SelectionKey.OP_WRITE);
        }
    }

    public void handleWrite(SelectionKey key) throws IOException {
        SocketChannel channel = (SocketChannel) key.channel();
        ByteBuffer buffer = (ByteBuffer) key.attachment();
        buffer.flip();
        channel.write(buffer);
        if (!buffer.hasRemaining()) {
            key.interestOps(SelectionKey.OP_READ);
        }
        buffer.compact();
    }
}
```

# IO models

Java NIO - non-blocking, sockets



# IO models

Java NIO - non-blocking, sockets

- Not thread-safe selection keys
- OP\_WRITE
- Old Sun Windows implementation problems:
  - OP\_WRITE | OP\_READ
  - Sockets not closed



# IO models

## Java NIO.2

- Start non-blocking I/O operations
- Provide notifications when I/O completes:
  - `java.util.concurrent.Future<V>` - pending result
  - `CompletionHandler` - complete result (callback)

# IO models


## Java NIO.2 - Threading

- Default Group
- Custom Groups
  - Fixed Thread Pool
  - Cached Thread Pool


# IO models

## Java NIO.2 - Threading


### Fixed thread pool



### Cached thread pool



### Default Group



# IO models

## Java NIO.2 - Issues

- Fixed Thread Pool
  - blocking completion handlers
- Cached Thread Pool
  - context switches
  - unbounded queuing needed, OOM
- ByteBuffer pool
- write() -> complete() -> write()

# Network architecture patterns

# Network architecture patterns

- Serial
- Process per connection
- Thread per connection
- Preforking
- Thread pool
- Evented (Reactor)
- Hybrids